

E. Install python

```
Python version ::
sys.version_info(major=3, minor=6, micro=5, releaselevel='final', serial=0)
=====
OpenCV version ::
3.4.0
=====
Tensorflow version ::
1.4.1
```

F. Install Emacs

I am joking, Emacs is always installed ;-).

III. DAY 2 [FRI JUL 6 19:21:03 CDT 2018]

9.520/6.860: Statistical Learning Theory and Applications, Fall 2017

Instructors: Tomaso Poggio (TP), Lorenzo Rosasco (LR), Georgios Evangelopoulos (GE)

A. Course description

The course covers foundations and recent advances of Machine Learning from the point of view of Statistical Learning and Regularization Theory.

Understanding intelligence and how to replicate it in machines is arguably one of the greatest problems in science. Learning, its principles and computational implementations, is at the very core of intelligence. During the last decade, for the first time, we have been able to develop artificial intelligence systems that can solve complex tasks, until recently the exclusive domain of biological organisms, such as computer vision, speech recognition or natural language understanding: cameras recognize faces, smart phones understand voice commands, smart speakers/assistants answer questions and cars can see and avoid obstacles.

The machine learning algorithms that are at the roots of these success stories are trained with labeled examples rather than programmed to solve a task. Among the approaches in modern machine learning, the course focuses on regularization techniques, that provide a theoretical foundation to high-dimensional supervised learning. Besides classic approaches such as Support Vector Machines, the course covers state of the art techniques using sparsity or data geometry (aka manifold learning), a variety of algorithms for supervised learning (batch and online), feature selection, structured prediction, and multitask learning and principles for designing or learning data representations. Concepts from optimization theory useful for machine learning are covered in some detail (first order methods, proximal/splitting techniques, ...).

The final part of the course will focus on deep learning networks. It will introduce an emerging theory formalizing three key areas for the rigorous characterization of deep learning: approximation theory – which functions can be represented efficiently?; optimization theory – how easy is it to minimize the training error?; and generalization properties – is classical learning theory sufficient for deep learning? It will also outline a theory of hierarchical architectures that aims to explain how to build machine that learn using cortex principles and similar to how children learn: from few labeled and many more unlabeled data.

The goal of the course is to provide students with the theoretical knowledge and the basic intuitions needed to use and develop effective machine learning solutions to challenging problems.

B. Prerequisites

We will make extensive use of basic notions of calculus, linear algebra and probability. The essentials are covered in class and in the math camp material. We will introduce a few concepts in functional/convex analysis and optimization. Note that this is an advanced graduate course and some exposure on introductory Machine Learning concepts or courses is expected. Students are also expected to have basic familiarity with MATLAB/Octave.

C. Grading

Pset and project tentative dates: (slides).

D. Problem Sets

Problem Set 1 Problem Set 2 Problem Set 3 Problem Set 4

Submission instructions: Follow the instructions included with the problem set. Use the latex template for the report (there is a maximum page limit). Submit your report online through stellar.mit by the due date/time and a printout in the first class after the due date.

E. Projects

Reports are 1–page, extended abstracts using NIPS style files.

1) *Projects archive*: List of Wikipedia entries, created or edited as part of projects during previous course offerings.

F. Syllabus

URL.

Follow the link for each class to find a detailed description, suggested readings, and class slides. Some of the later classes may be subject to reordering or rescheduling.

Class	Date	Title	Instructor(s)
Class 01	Wed Sep 06	The Course at a Glance	TP
Class 02	Mon Sep 11	The Learning Problem and Regularization	LR
Class 03	Wed Sep 13	Reproducing Kernel Hilbert Spaces	LR
Class 04	Mon Sep 18	Positive Definite Functions, Feature Maps and Mercer Theorem	LR
Class 05	Wed Sep 20	Tikhonov Regularization and the Representer Theorem	LR
Class 06	Mon Sep 25	Logistic Regression and Support Vector Machines	LR
Class 07	Wed Sep 27	Regularized Least Squares	LR
Class 08	Mon Oct 02	Iterative Regularization via Early Stopping	LR
Class 09	Wed Oct 04	Learning with Stochastic Gradients	LR
Class 10	Wed Oct 11	Large Scale Kernel Methods	LR
Class 11	Mon Oct 16	Sparsity Based Regularization	LR
Class 12	Wed Oct 18	Convex Relaxation and Proximal Gradient	LR
Class 13	Mon Oct 23	Structured Sparsity Regularization	LR
Class 14	Wed Oct 25	Multiple Kernel Learning	LR
Class 15	Mon Oct 30	Learning Theory	LR
Class 16	Wed Nov 01	Generalization Error and Stability	LR
Class 17	Mon Nov 06	Online Learning II	Sasha Rakhlin
Class 18	Wed Nov 08	Online Learning II	Sasha Rakhlin
Class 19	Mon Nov 13	Data Representation by Design	GE
Class 20	Wed Nov 15	Learning Data Representation: Dictionary Learning	GE
Class 21	Mon Nov 20	Learning Data Representation: Neural Networks	GE
Class 22	Wed Nov 22	Deep Learning Theory: Approximation	TP
Class 23	Mon Nov 27	Deep Learning Theory: Optimization	TP
Class 24	Wed Nov 29	Deep Learning Theory: Generalization	TP
Class 25	Mon Dec 04	Learning Data Representation: Invariance and Selectivity	TP
Class 26	Wed Dec 06	Deep Networks and Visual Cortex	TP
Class 27	Mon Dec 11	Poster presentations (2 sessions)	

G. Class 1. Course at a Glance

1) *Description*: We introduce and motivate the main theme of much of the course, setting the problem of supervised learning from examples as the ill-posed problem of approximating a multivariate function from sparse data. We present an overview of the theoretical part of the course and sketch the connection between classical Regularization Theory with its RKHS-based algorithms and Learning Theory. We briefly describe several different applications ranging from vision to computer graphics, to finance and neuroscience. The last third of the course will be on data representations for learning and deep learning. It will introduce recent theoretical developments towards a) understanding why deep learning works and b) a new phase in machine learning, beyond classical supervised learning: how to learn in an unsupervised way representations that significantly decrease the sample complexity of a supervised learning.

2) *Slides*:

- Slides for this lecture: PDF.

Youtube video class 2015.: [href6AWZS4Ho2Z8](https://www.youtube.com/watch?v=6AWZS4Ho2Z8)video

Link here

2017 Course – Center for Brains, Minds and Machines (CBMM): [hrefQ5itLKscYT](https://www.youtube.com/watch?v=Q5itLKscYT)video

Link here

3) *Relevant Reading*:

- Mnih et. al. (Deep Mind), Human-level control through deep reinforcement learning, *Nature* 518, pp. 529–533, 2015.
- Nature Insights, Machine Intelligence (with review article on Deep Learning), *Nature*, Vol. 521 No. 7553, pp. 435–482, 2015.